

Perceptron

Machine Learning Techniques

Dr. Ashish Tendulkar
IIT Madras

Perceptron is a binary classification algorithm.

We will cover

- Five different components of this algorithm, just like any other ML algorithm along with mathematical details.
- **Implementation** from scratch in Python with Numpy.

Perceptron was motivated from neurons

- Invented in 1958 by Frank Rosenblatt and was intended to be a machine, rather than a program.
- Perceptron was meant to be a rough model of how individual neurons work in the brain.

Let's look at the first component that is **training data**, which is very similar to other classification algorithms.

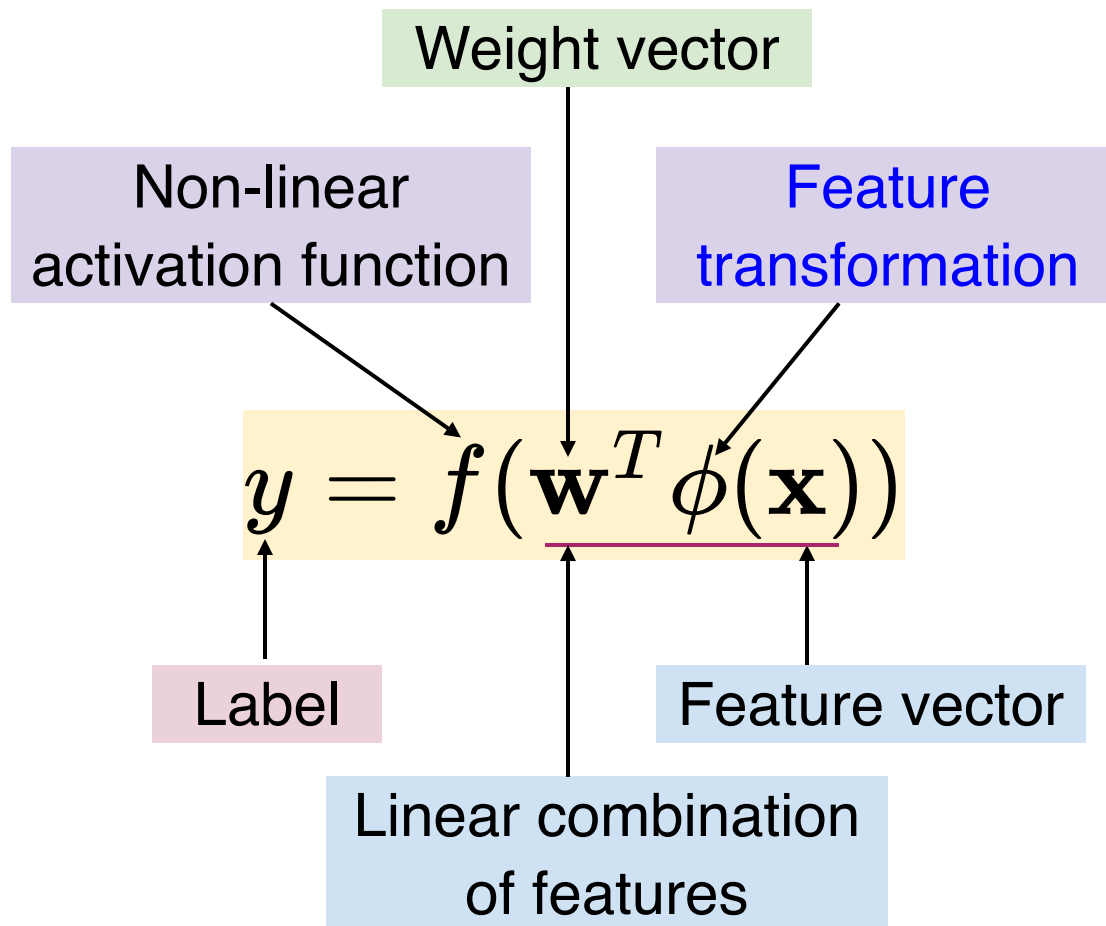
Training Data

- Feature matrix: $\mathbf{X}_{n \times m}$
- Label vector: $\mathbf{y}_{n \times 1}$

Note that perceptron can solve **only binary classification** problems. Hence $y^{(i)} \in \{-1, +1\}$

Let's look at the second component that is **model**, which is inspired from brain neurons.

Model $h_{\mathbf{w}}(\mathbf{x})$



Model $h_{\mathbf{w}}(\mathbf{x})$

$$y = f(\mathbf{w}^T \phi(\mathbf{x})) = f(z)$$

where, $f(\cdot)$ is a **non-linear activation function**. Here we use **sign or threshold function** as $f(\cdot)$:

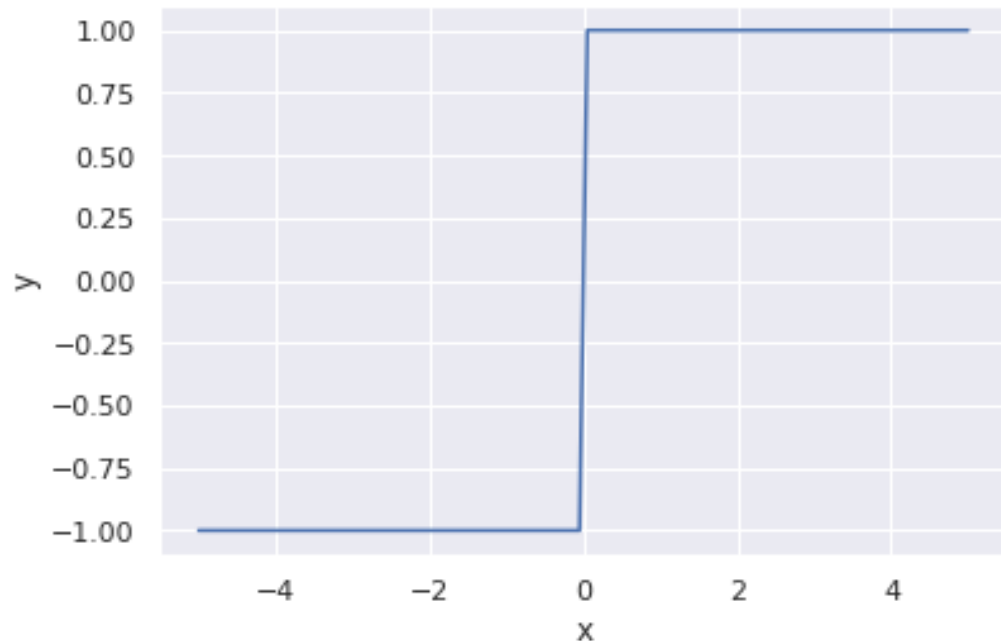
$$f(z) = \begin{cases} +1, & \text{if } z \geq 0 \\ -1, & \text{otherwise (i.e. } z < 0) \end{cases}$$

It can also be written as

$$y = \text{sign}(\mathbf{w}^T \phi(\mathbf{x}))$$

Model visualization

Remember $y = \text{sign}(\mathbf{w}^T \phi(\mathbf{x}))$



Note that for values of $x < 0$, we have $y = -1$

And for values of $x \geq 0$, we have $y = +1$

Now that we have seen the model of perceptron, let's look at the third component that is **loss function**.

Loss function

Let $\hat{y}^{(i)} \in \{-1, +1\}$ be the **prediction** from perceptron and $y^{(i)}$ be the **actual label** for i -th example. The **error** $e^{(i)}$ is calculated as

$$e^{(i)} = \begin{cases} 0, & \text{if } \hat{y}^{(i)} = y^{(i)} \\ -\mathbf{w}^T \phi(x^{(i)}) y^{(i)}, & \text{otherwise (i.e. } \hat{y}^{(i)} \neq y^{(i)}) \end{cases}$$

For **correctly classified examples**, the **error** is 0

For **misclassified examples**, the **error** is $-\mathbf{w}^T \phi(x^{(i)}) y^{(i)}$

Loss function

The error can be compactly written as:

$$\begin{aligned} e^{(i)} &= \max(0, -\mathbf{w}^T \phi(x^{(i)}) y^{(i)}) \\ &= \max(0, -h_{\mathbf{w}}(x^{(i)}) y^{(i)}) \end{aligned}$$

Loss function Illustration

$h_{\mathbf{w}}(\mathbf{x}^{(i)}) : y = \text{sign}(\mathbf{w}^T \phi(\mathbf{x}^{(i)}))$ is either +1 or -1.

If the **decision is correct**, then

$$h_{\mathbf{w}}(\mathbf{x}^{(i)}) = +1 \text{ and } y^{(i)} = +1 \quad \text{or :} \quad h_{\mathbf{w}}(\mathbf{x}^{(i)}) = -1 \text{ and } y^{(i)} = -1$$

$$\begin{aligned} e^{(i)} &= \max(0, -h_{\mathbf{w}}(\mathbf{x}^{(i)})y^{(i)}) \\ &= \max(0, -(1 \times 1)) \\ &= \max(0, -1) \\ &= 0 \end{aligned}$$

$$\begin{aligned} e^{(i)} &= \max(0, -h_{\mathbf{w}}(\mathbf{x}^{(i)})y^{(i)}) \\ &= \max(0, -(-1 \times -1)) \\ &= \max(0, -1) \\ &= 0 \end{aligned}$$

Note that in both the cases the error is 0.

Loss function Illustration

If the decision is wrong, then

$$h_{\mathbf{w}}(\mathbf{x}^{(i)}) = +1 \text{ and } y^{(i)} = -1$$

$$\begin{aligned} e^{(i)} &= \max(0, -h_{\mathbf{w}}(\mathbf{x}^{(i)})y^{(i)}) \\ &= \max(0, -(+1 \times -1)) \\ &= \max(0, 1) = 1 \end{aligned}$$

$$h_{\mathbf{w}}(\mathbf{x}^{(i)}) = -1 \text{ and } y^{(i)} = +1:$$

$$\begin{aligned} e^{(i)} &= \max(0, -h_{\mathbf{w}}(\mathbf{x}^{(i)})y^{(i)}) \\ &= \max(0, -(-1 \times +1)) \\ &= \max(0, 1) = 1 \end{aligned}$$

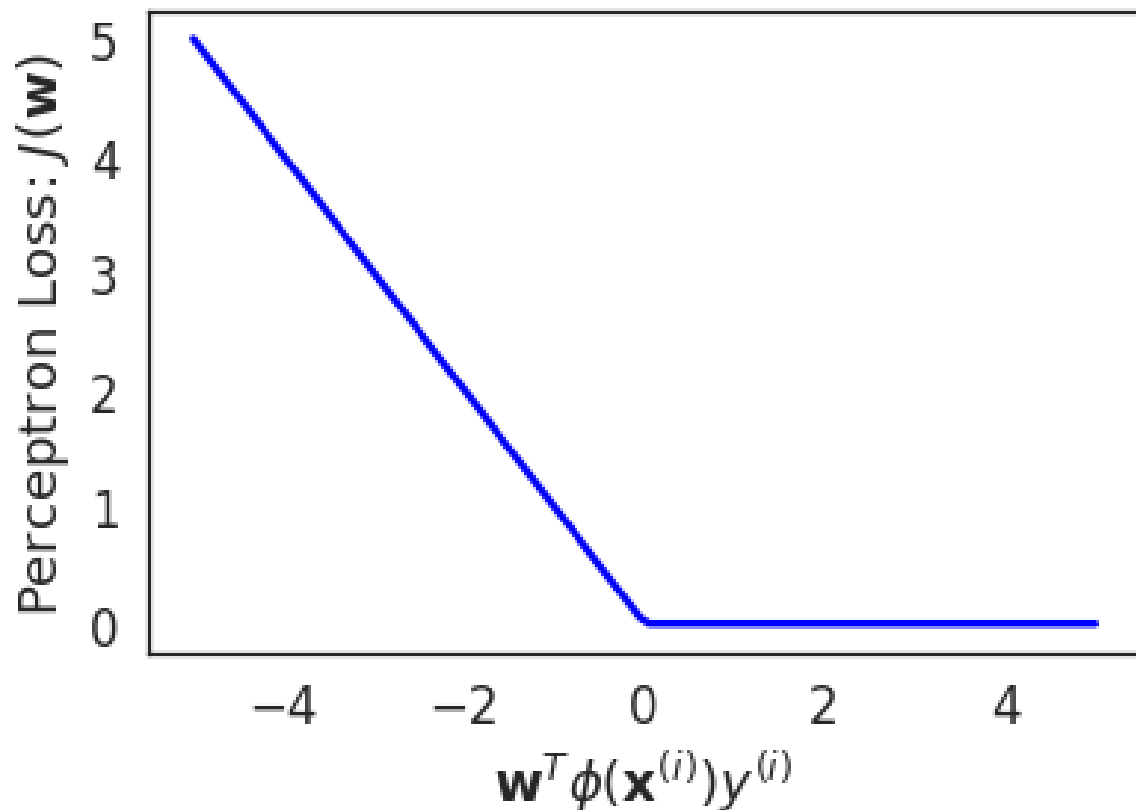
Note that in both the cases the error is 1.

Loss function

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^n e^{(i)} \\ &= \sum_{i=1}^n \max(0, -\mathbf{w}^T \phi(x^{(i)}) y^{(i)}) \\ &= \sum_{i=1}^n \max(0, -y^{(i)} h_{\mathbf{w}}(x^{(i)})) \end{aligned}$$

The error is a **piecewise linear function**: it is **zero** in the correctly classified regions and a **linear function of \mathbf{w}** in mis-classified region.

$J(\mathbf{w})$ is not differentiable in \mathbf{w} .



We can control the loss is by adjusting the value of \mathbf{w} .

The loss is directly proportional to \mathbf{w} .

Thus for **misclassified example**, we can **reduce loss by reducing \mathbf{w}** .

And for **correctly classified examples**, we leave \mathbf{w} **unchanged**.

The next task is to obtain the **weight vector** that **minimizes the loss**.

We will look at the optimization procedure used in perceptron. This procedure is known as **perceptron update rule**.

Optimization procedure

1. Initialize $\mathbf{w}^{(0)} = \mathbf{0}$

2. For each training example $(\mathbf{x}^{(i)}, y^{(i)})$:

$$\hat{y}^{(i)} = \text{sign}(\mathbf{w}^T \phi(\mathbf{x}^{(i)}))$$

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + \alpha (y^{(i)} - \hat{y}^{(i)}) \phi(\mathbf{x}^{(i)})$$

Note for **correctly classified examples**, $(y^{(i)} - \hat{y}^{(i)}) = 0$ and hence there is **no change in the weight vector**.

Linear separable examples lead to convergence of the algorithm with zero training loss, else it oscillates.

Let's understand **perceptron update rule** for various values of $\hat{y}^{(i)}$ and $y^{(i)}$:

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + \alpha (y^{(i)} - \hat{y}^{(i)}) \phi(\mathbf{x}^{(i)})$$

(Case 1: Correct classification) $\hat{y}^{(i)} = y^{(i)}$:

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + \alpha \times \mathbf{0} \times \phi(\mathbf{x}^{(i)}) = \mathbf{w}^{(t)} + \mathbf{0} = \mathbf{w}^{(t)}$$

(Case 2: Negative class misclassification) $y^{(i)} = -1$ and $\hat{y}^{(i)} = 1$:

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + \alpha \times (-1 - 1) \times \phi(\mathbf{x}^{(i)}) = \mathbf{w}^{(t)} - 2\alpha\phi(\mathbf{x}^{(i)})$$

(Case 3: Positive class misclassification) $y^{(i)} = 1$ and $\hat{y}^{(i)} = -1$:

$$\mathbf{w}^{(t+1)} := \mathbf{w}^{(t)} + \alpha \times (1 - (-1)) \times \phi(\mathbf{x}^{(i)}) = \mathbf{w}^{(t)} + 2\alpha\phi(\mathbf{x}^{(i)})$$

Optimization procedure: Convergence

- On linearly separable data, the optimization procedure eventually converges.

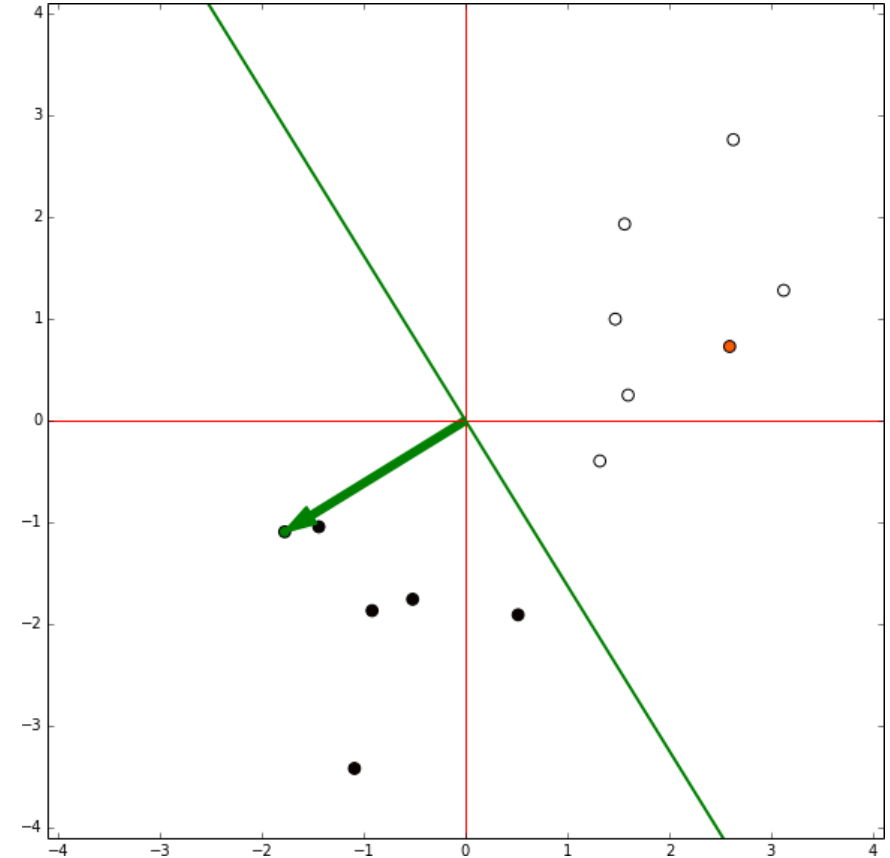
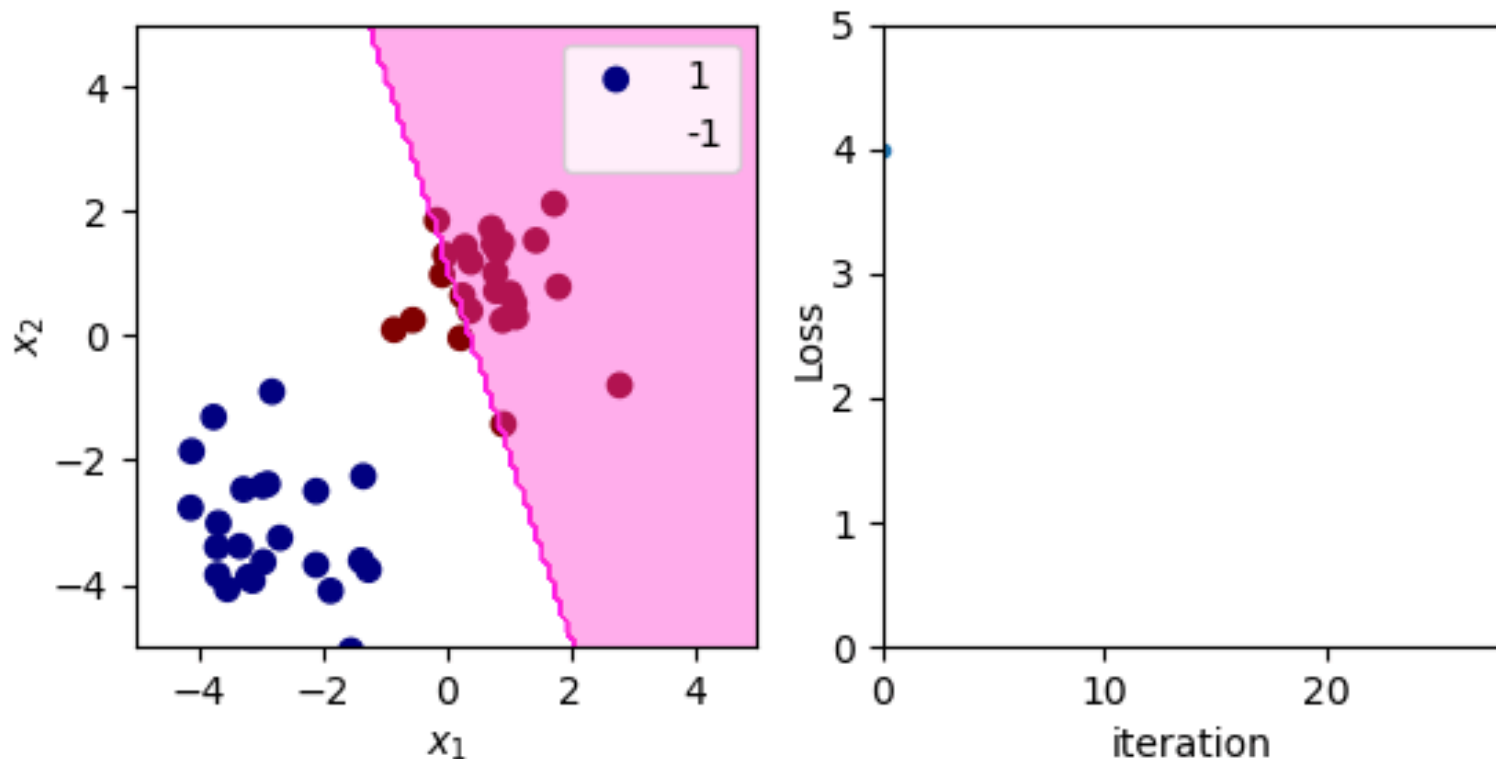


Image source: [Wikipedia.org](https://en.wikipedia.org/wiki/Support_vector_machine)

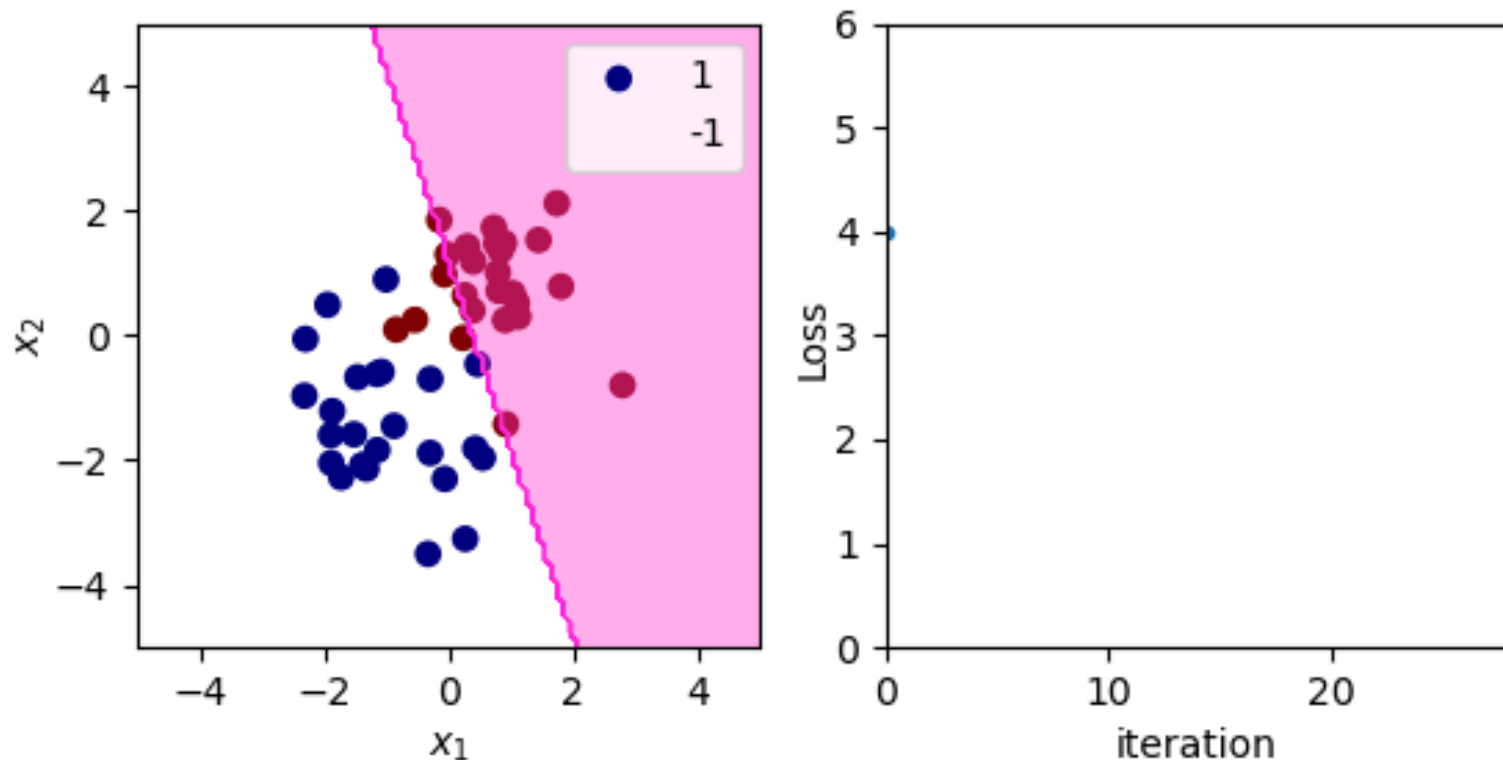
Optimization procedure: Convergence

- On linearly separable data, the optimization procedure converges.



Optimization procedure: Oscillations

- On **non** linearly separable data, the optimization procedure **never** converges.

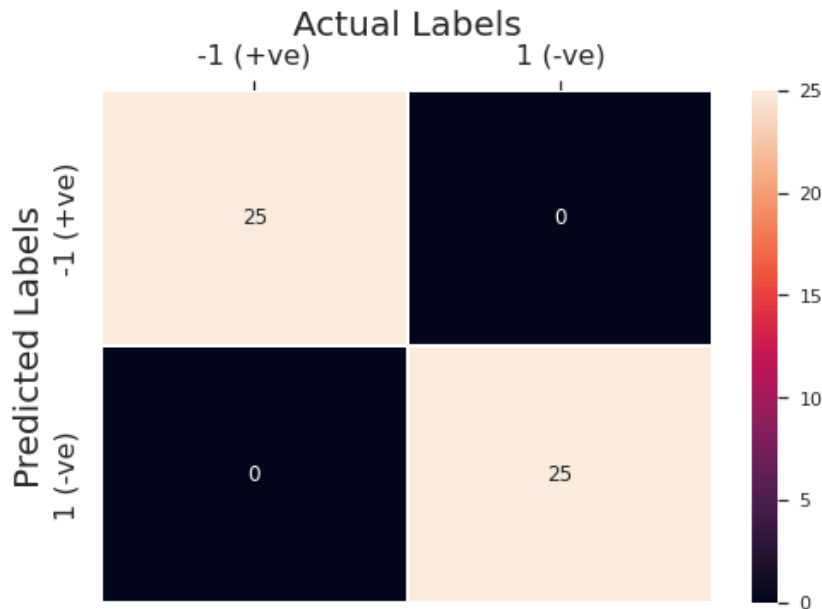


Finally let's look at the **evaluation metrics**.
They are same as other classification
algorithms.

Evaluation metrics

- Calculate **confusion matrix** based on **predicted and actual labels**.
- Calculate classification metrics like **precision, recall, accuracy, F1-score** from confusion matrix.

Confusion matrix and evaluation metrics: Example



$$\text{Precision (P)} = \frac{TP}{TP + FP} = \frac{25}{25 + 0} = 1.0$$

$$\text{Recall (R)} = \frac{TP}{TP + FN} = \frac{25}{25 + 0} = 1.0$$

$$\text{F1 Score} = \frac{2 \times P \times R}{P + R} = \frac{2 \times 1 \times 1}{1 + 1} = 1.0$$

- TP = 25
- FP = 0
- TN = 25
- FN = 0

Confusion matrix implementation

```
def compute_confusion_matrix(y, y_predicted):  
  
    # Create a 2D matrix of size n by n  
    confusion_matrix = np.zeros((n,n))  
  
    # Populate entries of confusion matrix  
    for i_x,i_y in zip(y_predicted, y):  
        confusion_matrix[i_x,i_y] +=1  
  
    return confusion_matrix
```